

0  
answers

### Q: Java Heap space error in Hadoop

I have setup **hadoop** cluster with single node as OS : CentOS Linux release 7.0.1406 **Hadoop** : 2.6.0  
Java : 1.7 My problem is with **hadoop heap** memory. I am currently reading file of 50 MB in reducer ...  
information. **Heap** Memory used 282.52 MB of 471 MB **Heap** Memory. Max **Heap** Memory is 8.89 GB.  
More Description For **Error**: What is missing in the configuration. Can anybody please help me?? ...

java linux hadoop memory

asked 13h ago by [Sanjay Bhosale](#)

0  
answers

### Q: Hadoop Error: Java heap space

So, after seeing the a percent or so of running the job I get an **error** that says, "**Error**: Java **heap** space"  
and then something along the lines of, "Application container killed" I am literally ... running an empty  
map and reduce job. However, the job does take in an input that is, roughly, about 100 gigs. For  
whatever reason, I run out of **heap** space. Although the job does nothing. I am using default ...

hadoop

asked may 7 by [Xibz](#)

0  
answers

### Q: Java heap space error while running hadoop

Every time I try to run a **hadoop** jar , I'm getting Java **heap** space **error**. I have installed **hadoop**-1.2.1 in  
my VMWare. This is my conf/mapred-site.xml: I have also added the following to **hadoop** ... -env.sh:  
Even after, each time I run the job it shows Java **heap** space exception and the job fails. Could some one  
help me out. Whenever I run my mapreduce my Linux becomes very slow. Too slow that it takes long  
even to open a browser. ...

java hadoop heap-memory

asked sep 6 by [NeethuPL](#)

1  
answer

### Q: Hadoop java mapper -copyFromLocal heap size error

As part of my Java mapper I have a command executes some code on the local node and copies a local  
output file to the **hadoop** fs. Unfortunately I'm getting the following output: **Error** occurred ... during  
initialization of VM Could not reserve enough space for object **heap** I've tried adjusting  
mapred.map.child.java.opts to -Xmx512M, but unfortunately no luck. When I ssh into the node, I can ...

java jvm hadoop heap mapreduce

asked sep 23 '11 by [Joris](#)

1  
answer

### Q: Heap error when using custom RecordReader with large file

Guide' but I get a **heap error** when trying to read in to the BytesWritable. I believe this is because the  
byte array is of size 85713669, but I'm not sure how to overcome this issue. Here is the code: } ... I've  
written a custom file reader to not split my input files as they are large gzipped files and I want my first  
mapper job to simply gunzip them. I followed the example in '**Hadoop** The Definitive ...

java hadoop heap

asked feb 11 '13 by [Shane](#)

0  
answers

### Q: Hadoop on local machine - Error occurred during initialization of VM, Could not reserve enou...

I'm trying to run Wikipedia Miner's CSV extraction on a wikipedia dump. I set up **Hadoop** 1.0.4 in single-  
node cluster configuration on a virtualized Mint 14 32-bit OS, using the instructions here ... . The works  
fine, but when I try to run the CSV extraction from Wikipedia Miner, I get a **error**. I'm not sure what  
information would be useful to provide. I allocated 4 cores (actually 2 hyperthreaded ...

hadoop linuxmint

asked mar 6 '13 by [feet](#)

0  
answers

### Q: Out of heap error when creating Index in Apache Hive

of RAM, 8 cores and 8 data disks. We use **Hadoop** v 2.4.1 and Hive v 0.13.1 . Before we got this far we  
had problems with running out of **heap** in the hive "console" but after increasing the max size we ...  
instead run into this new problem. We have tried various **heap** sizes (and garbage collectors/settings) but  
it seems to not make any difference (or at least not make the job complete without **errors** ...

indexing out-of-memory hive

asked aug 26 by [Magnus Eriksson](#)

2  
answers

### Q: Hadoop 2.4 maven build OutOfMemoryError: Java heap space

I was trying to build the complete source code of **Hadoop** on my mac. When I try building I constantly  
run into the **error**: "java.lang.OutOfMemoryError: Java **heap** space". I tried to set the maven opts ...

maven hadoop heap-memory

asked nov 26 by [Sachin Ananth Navale](#)

0  
answers

### Q: Heap space issue with hadoop-streaming using hadoop 2x

I am having a serious issue with **hadoop** 2.4. i am trying to run my use case which is a groovy script  
using **hadoop**-streaming jar. It results in **error**: Java **heap** space - No memory Here is my ...  
configuration info yarn-site.xml mapred-site.xml I have 8GB RAM total in every node and there is no  
reduce task. My command line to run the job is as follows:- \$HADOOP\_INSTALL/bin/**hadoop** jar ...

hadoop heap-memory hadoop-streaming

asked jul 9 by [user3721344](#)

1  
answer

### Q: hadoop mapper over consumption of memory(heap)

and with the mapper code I dont expect my **heap** memory to ever reach 256M, yet it fails with java **heap**  
space **error**. I will be thankful if you can give some insight into why the mapper is consuming so much  
memory. EDIT: ... I wrote a simple hash join program in **hadoop** map reduce. The idea is the following: A  
small table is distributed to every mapper using DistributedCache provided by **hadoop** framework. The  
large table ...

hadoop mapreduce hashmap mapper

asked mar 9 '13 by [mc\\_87](#)

2

answers

### Q: CDH 4.1: Error running child : java.lang.OutOfMemoryError: Java heap space

I've trying to overcome sudden problem. Befor that problem I've used old VM. I've downloaded the new one VM and still can't make my job run. I get Java **heap space error**. I've already read this one ... post: out of Memory **Error in Hadoop** Here is my configs from /etc/**hadoop**/conf: sudo vi **hadoop-env.sh** Here is my mapred-site.xml Nothing helps :( Here is the console output: Here is my log: What can I try next? Please help. Thank you. ...

hadoop mapreduce cloudera yarn

asked dec 2 '12 by Sergey

1

answer

### Q: increase jvm heap space while running from hadoop unix

I am running a java class test.java from **hadoop** command : I am using a string builder, and its size is going out of memory : I know in java we can run a java program by providing a **heap space** ... size : How can I do this while running with **hadoop**, if I run it like : it throws exception : I would like to change the **heap** size permanently. ...

java unix exception hadoop

asked aug 31 '13 by naveen kumar

2

answers

### Q: Error: Java heap space

In Ubuntu, when I am running the **hadoop** example : In log, I am getting the **error** as : INFO mapred.JobClient: Task Id : attempt\_201303251213\_0012\_m\_000000\_2, Status : FAILED **Error**: Java ... **heap space** 13/03/25 15:03:43 INFO mapred.JobClient: Task Id :attempt\_201303251213\_0012\_m\_000001\_2, Status : FAILED **Error**: Java **heap space**13/03/25 15:04:28 INFO mapred.JobClient: Job Failed ...

hadoop

asked mar 25 '13 by Senthil Porunan

2

answers

### Q: hadoop cannot allocate memory java.io.IOException: error=12

i am getting the following **error** on **hadoop** greenplum the server has 7G ram and 1G swap. **heap** size is 1024m and mapred.child.opts is set to 512m. any ideas? ...

java linux hadoop

asked sep 27 '13 by krisdigitx

2

answers

### Q: Error in running Hadoop example on single node cluster

I'm running the word count example provided by **hadoop**. The following is the **error** that i am getting: Information about my installation: I followed the installation instructions from Michael ... Noll's link Here The only change is that i haven't created a separate **hadoop** user. I'm using it under my login. ...

hadoop mapreduce distributed-computing

asked jul 29 '11 by Chander Shivasani

1

answer

### Q: hadoop cannot start slave "fatal error"

i am starting my **hadoop** with 4 slaves and all works fine except for one machine. I have created them the exact same way. the **error** i receive when running ./start-all.sh is: the details of the log ... : the details of the third log is: im not sure of the problem could someone help me out? I have created the slaves the exact same way. **heap** is set to 1GB. thank you so much ! ...

java hadoop startup fatal-error

asked mar 14 '13 by cs\_newbie

4

answers

### Q: Hadoop: Error in configuring object

details provided at: <http://www.michael-noll.com/blog/2011/04/09/benchmarking-and-stress-testing-an-hadoop-cluster-with-terasort-testdfsio-nnbench-mrbench/> ...

java exception hadoop mapreduce cloudera

asked jan 24 '12 by Chander Shivasani

0

answers

### Q: How to specify Heapsize in an EMR streaming job

I have a **hadoop** streaming job running (and failing) on amazon EC2. My job is as follows: the job fails due to java **heap size error** After looking around, I understand **heap** space allocated ... is not enough. But I couldn't find how to do specify the **heap** size for a **hadoop** streaming job on EC2 instance. Environment variables such as HADOOP\_OPTS doesn't exist. How can I fix this? Thanks! ...

python hadoop amazon-web-services heap-size

asked apr 21 by user3394040

0

answers

### Q: hadoop task node can not launch because of small heap size? solved. it is because of hdfs-si...

from **hadoop** package, I get following **error**, then I go to task node and check log files, which is in logs/userlogs/attempt\_201203152028\_0002\_r\_000001\_1/stdout I do the following, Then I check ... I have configured the **hadoop** cluster with 8 slave nodes, I can launch the namenode and datanode smoothly, and I can save files in hdfs, however, I want to test the program using example java program ...

hadoop size heap

asked mar 15 '12 by user974270

1

answer

### Q: hadoop windows (work ok) linux java heap space

) (using cygwin) on a 4 GB RAM machine, the application works fine, but when I run it on linux- ubuntu on a 2 GB RAM machine, it process some rows but then it throws a "Java **heap space**" **error**, or sometimes ... the thread is killed. For the linux: I already tried to change the **hadoop** export HEAP\_SIZE and also the Xmx and Xms parameters on the app and it made some difference but not too much, the **error** ...

java linux heap hadoop space

asked nov 1 '10 by KoRnE

0

answers

### Q: running fpg algorithm of mahout on hadoop as cluster mod

and HADOOP\_HOME and MAHOUT\_LOCAL , .... in .bashrc file in the user's environment **hadoop** command: **hadoop**@master mahout\$ bin/mahout fpg -i /home/**hadoop**/output.dat -o patterns -method

mapreduce -k 50 -s 2 **error**: Please help me. I tried but could not fix the **error**. ... I install mahout-0.7 and **hadoop**-1.2.1 on linux (centos).**hadoop** config as multi\_node. I created a user named **hadoop** and install mahout and **hadoop** in path /home/**hadoop**/opt/ I set MAHOU\_HOME ...

[hadoop](#) [mahout](#)

asked 2d ago by [raha\\_gh](#)

1  
answer

### Q: PIG using HCatLoader, Java heap space error

- Connected to metastore. 2014-04-11 01:12:37,785 [main] **ERROR** org.apache.pig.tools.grunt.Grunt - **ERROR** 2998: Unhandled internal **error**. Java **heap** space Details at logfile: /home/aphadoop/pig\_1397158893348.log Does anybody encountered the same **error** and solved, Please let me know. ... I am using hive-0.12.0,pig-0.12.0,mysql-5.6 and **hadoop**-1.2.1 in pseudo distribution mode. I configured PIG\_CLASSPATH etc.. details according to the link, <https://cwiki.apache.org/confluence/display/...>

[hadoop](#) [apache-pig](#) [hcatalog](#)

asked apr 12 by [Dorababu G](#)

0  
answers

### Q: Error running child : java.lang.OutOfMemoryError: Java heap space

I have read a lot on the internet, but found no solution for my problem. I use **Hadoop** 2.6.0. The main goal for the MapReduce is to run through a SequenceFile and do some analysis on the key/value ... pairs. Here the output from STDOUT My configurations are nearly default, nothing related to the Java **heap** size. I've also tried this, which does not made a difference. The main programm ...

[java](#) [hadoop](#)

asked 13h ago by [Christian D.](#)

0  
answers

### Q: Java Heap space error in mahout

I have gone through the previous questions mentioned and tried a solution mentioned there. It didn't work out for me . It's the same Java **Heap** space **error** when running converting directory ... of **heap** space. For that I ran jps before running mahout job. This did not give me process id or anything as mentioned here. Maybe I am not doing it right. Please suggest ...

[java](#) [hadoop](#) [mapreduce](#) [mahout](#)

asked dec 7 by [user3527975](#)

3  
answers

### Q: After starting hadoop, i cannot start mahout ..!

When Starting Mahout i got an **error** message like this root@fazil-VPCEB45FG:/usr/local/mahout/bin# ./mahout **hadoop** binary is not in PATH,HADOOP\_HOME/bin,HADOOP\_PREFIX/bin, running locally ... **Error** occurred during initialization of VM Could not reserve enough space for object **heap** Could not create the Java virtual machine. i have installed openjdk0.6, Whether openjdk is supported by mahout ...

[mahout](#)

asked apr 13 '13 by [user2260984](#)

0  
answers

### Q: Hadoop error in shuffle in fetcher#1

It failed after 11 maps tasks out of 200 planned. I'm running with **Hadoop** 0.22.0 I'm getting the following **error**: org.apache.hadoop.mapreduce.task.reduce.Shuffle\$ShuffleError: **error** in shuffle ... I'm running a parsing job in **hadoop**, the source is a 11GB map file with about 900,000 binary records each representing an HTML file, the map extract links and write them to the context. I have ...

[hadoop](#) [mapreduce](#)

asked nov 6 '13 by [Noam Hasson](#)

1  
answer

### Q: hadoop mapreduce gives Child Error

I am working with **hadoop** 1.2.1 on ubuntu 13.10. I am running the sort problem with the input file size of 25GB. But I am getting **error**: Sort example works perfectly fine with 10GB input file. I ...

[java](#) [hadoop](#)

asked sep 29 by [user2805242](#)

1  
answer

### Q: Hadoop: Can you silently discard a failed map task?

I am processing large amounts of data using **hadoop** MapReduce. The problem is that, occasionally, a corrupt file causes Map task to throw a java **heap** space **error** or something similar. It would be nice ... , if possible, to just discard whatever that map task was doing, kill it, and move on with the job, never mind the lost data. I don't want the whole M/R job to fail because of that. Is this possible in **hadoop** and how? ...

[java](#) [hadoop](#) [mapreduce](#)

asked jan 9 '14 by [miljanm](#)

1  
answer

### Q: Cloudera Hadoop MapReduce error: Task process exit with nonzero status of 65

it's status within 600 second what is regulated by mapred.task.timeout setting in mapred-site.xml. I'm not sure how to change this setting inside Cloudera 4.1.4 distribution of **Hadoop**? Also in following **error** log you can see some other **errors**: ...

[hadoop](#) [mapreduce](#) [hbase](#) [pentaho](#) [cloudera](#)

asked sep 20 '13 by [hakaa](#)

0  
answers

### Q: Querstion regarding hadoop-env.sh

I am Facing **Error**: Java **heap** space and **Error**: GC overhead limit exceeded So i started looking into **hadoop**-env.sh. so thats what i understand so far, Please correct me if i am wrong ... of 9GB But tasktracker is invoked with 7GB memory, so this wil conflict. as max memory for tasktracker and child JVMs invoked by tasktracker is 7GB, but they are consuming 9G. So the **heap** space **error** occured, is my calculation correct? ...

[java](#) [hadoop](#) [hadoop-streaming](#) [hadoop-partitioning](#) [hadoop2](#)

asked jun 27 by [user2950086](#)

2  
answers

### Q: java.lang.OutOfMemoryError on running Hadoop job

tried the solution of assigning more **heap** as mentioned in this thread: out of Memory **Error** in **Hadoop** WordCountMapper code: ... to the WordCount example provided with **Hadoop**. I've 4 classes in all to

carry out the processing: StanfordLemmatizer [contains goodies for lemmatizing from Stanford's coreNLP package v3.3.0], WordCount ...

java hadoop stanford-nlp

asked nov 27 '13 by Aditya

1  
answer

### Q: how to change mapper memory requirement in hadoop?

that lowering `mapred.max.split.size` was a good idea: more mappers with lower memory requirements. BUT, I got the "java.lang.OutOfMemoryError: Java **heap** space" **error** again, again and again. It seems that, i did not understand how **hadoop** works. Any suggestions? ... In a map-reduce job, i got the **error** "java.lang.OutOfMemoryError: Java **heap** space". Since I get this **error** in a mapper function; I thought that when I lower the input size to the mapper I will have ...

hadoop mapreduce

asked sep 11 '13 by ndemir

2  
answers

### Q: Hadoop conf to determine num map tasks

a Java **Heap Space error**. I've tried setting many different conf properties in my **Hadoop** cluster to make the job split into more tasks but nothing seems to have any effect. I have tried setting ... I have a job, like all my **Hadoop** jobs, it seems to have a total of 2 map tasks when running from what I can see in the **Hadoop** interface. However, this means it is loading so much data that I get ...

hadoop configuration mapreduce hdfs

asked jul 23 '13 by Katie

2  
answers

### Q: Mahout on Elastic MapReduce: Java Heap Space

I'm running Mahout 0.6 from the command line on an Amazon Elastic MapReduce cluster trying to canopy-cluster ~1500 short documents, and the jobs keep failing with a "**Error: Java heap space**" message ... Based on previous questions here and elsewhere, I've cranked up every memory knob I can find: `conf/hadoop-env.sh`: setting all the **heap** spaces there up to 1.5GB on small instances and even 4GB ...

hadoop heap mahout emr

asked apr 29 '12 by David M.

1  
answer

### Q: Hadoop Streaming Memory Usage

I'm wondering where the memory is used in the following job: **Hadoop** Mapper/Reducer **Heap Size**: Streaming API: Mapper: Reducer: Input File is a 350MByte file containing a single line full ... getting always out of memory **errors**. Is this a possible Bug in **Hadoop** (e.g. unaccary copies) or do I just don't understand some required memory intensive steps? Would be really thankful for any further hints. ...

java memory hadoop mapreduce

asked jul 31 '13 by mt\_

### A: Java Heap Space Error in running mahout item similarity job on Amazon EMR

/TaskConfiguration\_H2.html). Before making further changes you will need to determine what process is actually suffering from the **heap error** and then tune accordingly. If it does turn out ... to be container/task jvm **heap** then it can be adjusted with `configure-hadoop` ([http://docs.aws.amazon.com/ElasticMapReduce/latest/DeveloperGuide/emr-plan-bootstrap.html#PredefinedbootstrapActions\\_ConfigureHadoop](http://docs.aws.amazon.com/ElasticMapReduce/latest/DeveloperGuide/emr-plan-bootstrap.html#PredefinedbootstrapActions_ConfigureHadoop)). Also drop the memory-intensive bootstrap action, it is only for AMI 1.x. ...

answered jan 6 by ChristopherB

0  
answers

### Q: hadoop wordcount example error: Task process exit with nonzero status of 1

**hadoop-examples-1.2.1.jar** and got the stunning **error**: After getting this **error** I executed `sqoop import` from SQL server 2008 R2 in which I am getting same **error**. My job's log file also has same **error** ... I also verified by web UI that my cluster's **heap** size is 200 MB. So, to me memory doesn't seem problematic also userlogs are also there is `hadoop.temp.dir (/var/lib/hadoop)` and `stderr` file of logs says: Please let me know what should I do to resolve this issue. ...

hadoop mapreduce

asked feb 20 '14 by tom

2  
answers

### Q: Hadoop Map Reduce wordcount Shuffle Error: Exceeded MAX\_FAILED\_UNIQUE\_FETCHES; bailing-out

I have installed and configure **hadoop** as single node using manual from following site. <http://www.michael-noll.com/tutorials/running-hadoop-on-ubuntu-linux-single-node-cluster/#running-a-mapreduce-job> I have compiled wordcount example and run it but it takes long time and generate Any Clue? ...

java hadoop

asked jul 3 '12 by Suvendu

2  
answers

### Q: installing hadoop 2.3 on mac osx

i used the following tutorials to install **hadoop** on macos: **Hadoop** Tutorial 1, **Hadoop** Tutorial 2 when i try to format the namenode i get the following **error**: i wonder, because i have set my java home to: here also my `~/profile` file: and my **hadoop-env.sh**: my java version is: ...

java osx hadoop java-home

asked apr 3 by Gerald

0  
answers

### Q: Problems tuning memory consumption and containers number in Hadoop

)) Number of Containers: 8 ->  $\text{MIN}(2 * \text{CORE}, \text{Total Available RAM} / \text{MIN\_CONTAINER\_SIZE})$  233.6 This config, and some variants of the same, didn't work, and I get **error** like that Or also I get Java **Heaps** ... I'm trying to configure a **Hadoop** cluster. The default config work fine, but the resources consumption is poor. For example, the CPU total use in the slaves is slightly above 40%. I want to use ...

java xml hadoop memory-management

asked dec 10 by Tuxman

### Q: Writing a Hadoop Reducer which writes to a Stream

0  
answers

I have a **Hadoop** reducer which throws **heap** space **errors** while trying to produce very long output records. Is there a way to write a Reducer to use Streams for output, so that I can run through the data for the record without marshalling the whole record into memory? ...

hadoop streaming reducers

asked sep 10 by [mcintyre321](#)

2  
answers

### Q: Hadoop Failed to set permissions of path: \tmp\

I'm running (or obviously trying to) **Hadoop** 1.2.1 on my Windows machine inside cygwin. Unfortunately there is something terribly wrong with my **Hadoop**. I'm getting the following **error** when I'm trying ... to execute simple Pig script on local mode. I assume that there is something wrong with the **hadoop** installation or configuration files, but I'm new to **Hadoop** so is just an assumption. Can someone ...

windows hadoop cygwin

asked dec 3 '13 by [Anton Belev](#)

1  
answer

### Q: Getting java heap space error while running a mapreduce code for large dataset

-500MB. But for datasets above 1GB I am getting an **error** like this: I think my program is consuming too much memory and need to be optimized. I even tried to solve this by increasing my java **heap** space ... I am a beginner of MapReduce programming and have coded the following Java program for running in a **Hadoop** cluster comprising 1 NameNode and 3 DatanNodes : The program is run on a dataset like ...

java hadoop mapreduce

asked apr 13 by [Monami Sen](#)

1  
answer

### Q: What does this error when I'm trying to run an example in Apache Mahout?

I configured and ran examples with mahout but I get the following **error** : **hadoop** binary is not in PATH,HADOOP\_HOME/bin,HADOOP\_PREFIX/bin, running locally **Error** occurred during initialization ... of VM In my, , I defined the following aliases :  
JAVA\_HOME,HADOOP\_CONF\_DIR,MAHOUT\_CONF\_DIR,HADOOP\_HOME I already have configured this in : However, I have the following **error** : ...

hadoop config mahout

asked jul 26 '12 by [n99o](#)

2  
answers

### Q: Error running hadoop(Cloudera-2.0.0-cdh4.4.0) job from eclipse?

Hi i am running the **hadoop** wordcount example from eclipse and i am getting the following **error**:- I have added all the required JAR's in my project. Following is the code i am running:- Following are the jar files added to the project:- ...

java hadoop cloudera

asked nov 24 '13 by [Dennis](#)

2  
answers

### Q: Java Out of Memory exception in Ubuntu when using Flume/Hadoop

I'm getting an out of memory exception due to lack of Java **heap** space when I try and download tweets using Flume and pipe them into **Hadoop**. I have set the **heap** space currently to 4GB in the mapred ... - site.xml of **Hadoop**, like so: I am hoping to download tweets continually for two days but can't get past 45 minutes without **errors**. Since I do have the disk space to hold all of this, I am assuming ...

java hadoop heap flume

asked jul 29 '13 by [Andrew Martin](#)

3  
answers

### Q: Reducer's Heap out of memory

on occasion I also get **errors** about bash failing to get memory for what I assume is the spill operation. Would this be the **Hadoop** node running out of memory? If so would just turning down the **heap** ... So I have a few Pig scripts that keep dying in there reduce phase of the job with the **errors** that the Java **heap** keeps running out of space. To this date my only solution has been to increase Reducer ...

hadoop mapreduce apache-pig piglatin

asked jan 2 '12 by [NerdyNick](#)

1  
answer

### Q: Mahout on EMR Error: Java heap space

I've ran a clustering job on EMR. The dataset is huge. Everything worked well untill: So. The basic question is how to fix that?

java hadoop amazon-web-services mahout emr

asked sep 29 '12 by [denys](#)

1  
answer

### Q: Memory problems with Java in the context of Hadoop

I want to compute a multiway join in **Hadoop** framework. When the records of each relation get bigger from a threshold and beyond I face two memory problems, 1) **Error**: GC overhead limit exceeded, 2 ... ) **Error**: Java **heap** space. The threshold is the 1.000.000 / relation for a chain join and a star join. In the join computation I use some hash tables i.e. These **errors** occur when I hash the input ...

java hadoop garbage-collection heap

asked sep 22 '13 by [vpap](#)

2  
answers

### Q: A join operation using Hadoop MapReduce

a HashMap and then take a cross product. (eg. Join of two datasets in Mapreduce/**Hadoop**) This solution is very good and works for majority of the cases but in my case my issue is rather different. I am ... dealing with a data which has got billions of records and taking a cross product of two sets is impossible because in many cases the hashmap will end up having few million objects. So I encounter a **Heap** ...

hadoop mapreduce elastic-map-reduce

asked may 19 '13 by [Eastern Monk](#)