

0  
answers

### Q: Building Inverted Index exceed the Java Heap Size

the **error**. Which means that as the data increase number of items/product\_ids that I index for words like New or old etc get bigger which cause the **Heap** Size to overflow. So, the question is how can avoid java **heap** size overflow and accomplish this task. ... of data from large system). The building of inverted index get executed as a map reduce job on **Hadoop**. Inverted index is build with the help of scala. Structure of the inverted index as follows ...

scala   hadoop   avro   scalding

asked jul 31 '13 by [Null-Hypothesis](#)

0  
answers

### Q: Mahout Canopy Clustering, K-means Clustering : Java Heap Space - out of memory

centers by tweaking canopy parameters (T1, T2), it works very well. More then a certain number of canopy centers, the jobs keep failing with a "**Error: Java heap space**" message at 67% of reduce phase. K ... can find: **hadoop-env.sh**: setting all the **heap** spaces there up to 16GB on namenode and even 8GB on datanodes. **mapred-site.xml**: adding **mapred.{map, reduce}.child.java.opts** properties, and setting ...

hadoop   heap   mahout

asked jun 12 '13 by [Nebulax](#)

1  
answer

### Q: Sqoop incremental import to S3 Wrong FS error

When using the flag in the , the job will fail. Here is the full command: The exact same command without the `--check-column` work without any **error** and upload to the correct s3 directory ...

hadoop   cloudera   sqoop

asked feb 22 '13 by [ender](#)

1  
answer

### Q: CDH5.2: MR, Unable to initialize any output collector

-mapred-1.7.6-cdh5.2.0-hadoop2.jar, **hadoop-common-2.5.0-cdh5.2.0.jar**, **hadoop-mapreduce-client-core-2.5.0-cdh5.2.0.jar** I've run the following job I get the following **error**, is it a Java **heap** problem, any comments ? Thank you in advance ...

java   hadoop   mapreduce   yarn   cloudera-cdh

asked nov 14 by [guzu92](#)

0  
answers

### Q: Error when trying to run algorithm in Spark

based on cdh4 (<http://www.apache.org/dyn/closer.cgi/spark/spark-1.2.0/spark-1.2.0-bin-cdh4.tgz>) For the moment however I haven't been able to successfully run it as I keep getting a java **heap error** ... . The input file itself is only 136kb, so I don't think memory is the case. Below is an the command that I ran and the **error** message that I get. Command: Verbose output: Stderr from one of executor (executorID 0. Both of the executor had the same stderr output) ...

apache-spark

asked jan 15 by [Hafidz Zulkifli](#)

1  
answer

### Q: How to handle unsplittable 500 MB+ input files in hadoop?

out of memory **errors**, even with a java **heap** size of 6 GB. Now I wonder how I could split the data so that it better matches hadoop's 64 MB block size. I cannot simply split the big files ... I am writing a **hadoop** MapReduce job that is running over all source code files of a complete Debian mirror (~ 40 GB). Since the Debian mirror data is on a separate machine and not in the **hadoop** ...

hadoop   mapreduce

asked may 25 by [Michael](#)

1  
answer

### Q: Hadoop UniqValueCount Map and Aggregate Reducer for Large Dataset (1 billion records)

, but keep getting "Out of Memory" and Java **heap** size **errors** on **Hadoop** - at the same time, I am able to run this fairly easily on a single box using a Python Set (hashtable, if you will.) I am using a fairly ... I have a data set that has approximately 1 billion data points. There are about 46 million unique data points I want to extract from this. I want to use **Hadoop** to extract the unique values ...

hadoop   mapreduce   hadoop-streaming   elastic-map-reduce

asked jan 18 '13 by [Suman](#)

3  
answers

### Q: spark java.lang.OutOfMemoryError: Java heap space

my cluster: 1 master, 11 slaves, each has 6G memory. my setting: **spark.executor.memory**=4g, **Dspark.akka.frameSize**=512 here is my problem: First,I read some data(2.19G) from hdfs to ...

out-of-memory   apache-spark

asked jan 15 '14 by [hequn8128](#)

1  
answer

### Q: hadoop connection refused on port 9000

I want to setup a **hadoop**-cluster in pseudo-distributed mode for development. Trying to start the **hadoop** cluster fails due to refused connection on port 9000. These are my configs (pretty standard ... files: **hadoop-pschmidt-datanode-zaphod.log** **hadoop-pschmidt-namenode-zaphod.log** After reformatting the hdfs following output is displayed: Using : Using returns nothing, assuming ...

java   ubuntu   hadoop   hdfs

asked aug 19 '13 by [dyster fyr](#)

0  
answers

### Q: Hive Issue - java.lang.OutOfMemoryError: Java heap space

:597) at org.apache.hadoop.util.RunJar.main(RunJar.java:208) FAILED: Execution **Error**, return code -101 from org.apache.hadoop.hive ql.exec.DDLTask My Environment variables as defined in **hadoop-env.sh** are Definition of **hive-env.sh** Where could the problem be? Please advise. ... **heap** space at org.apache.thrift.protocol.TBinaryProtocol.readStringBody(TBinaryProtocol.java:353) at org.apache.thrift.protocol.TBinaryProtocol.readMessageBegin(TBinaryProtocol.java:215 ...

hadoop   hive

asked apr 22 by [user3528338](#)

### Q: OutOfMemoryError while Logistic regression in SparkR

1  
answer

269369 rows. The **error** message I got: Dimension of data (sample) : I have also tried hosting the same csv file over local file system as well as on HDFS. I think it will need more **Hadoop** ... datanotes to store large datasets?. If it is then how can I set up Spark **Hadoop** cluster to get rid of this. (or I am doing something wrong) Tip: I think increasing Java and Spark **heap** space will help me ...

r apache-spark

asked oct 6 by [Hanry](#)

#### A: libssh: Memory lost in simple implementation

It happened as I suspected. This could be the bug in . I ran my program in and valgrind showed no leaks. Here is the snippet of the output: So, of course, this is not a bug in .

answered jan 8 by [Rohit](#)

#### A: Hadoop example job fails in Standalone mode with: "Unable to load native-hadoop library"

issue is the you get. Check you input, increase the **heap** size, if necessary. You might also have a look at this related question: out of Memory **Error** in **Hadoop** ... The warning tells you that the compression codec is not (properly) installed for **Hadoop**. To install the compression, have a look at: <http://code.google.com/p/hadoop-snappy/> However, a more serious ...

answered dec 1 '12 by [Lorand Bendig](#)

0  
answers

#### Q: NegativeArraySizeException when creating a SequenceFile with large (>1GB) BytesWritable valu...

I have tried different ways to create a large **Hadoop** SequenceFile with simply one short(<100bytes) key but one large (>1GB) value (BytesWritable). The following sample works for out-of-box: <https://svn.apache.org/repos/asf/hadoop/common/trunk/hadoop-mapreduce-project/hadoop-mapreduce-client/hadoop-mapreduce-client-jobclient/src/test/java/org/apache/hadoop/mapred/BigMapOutput.java> which writes ...

hadoop out-of-memory heap large-files sequencefile

asked jun 9 by [user815613](#)

1  
answer

#### Q: Hadoop 2.4 cannot launch job on aws s3n

I deployed **Hadoop** 2.4 on AWS EC2 using S3 native file system as a replacement of HDFS. I tried several example apps, all gave me the following stack tracing msgs (an older thread on Jul 24 hang ... (core-site.xml): <https://wiki.apache.org/hadoop/AmazonS3> core-site.xml: I tried **Hadoop** v1 as well and turns out Hadoop1 on s3n file system works. But it seems it doesn't work for **Hadoop** v2. Please help. Thanks in advance. ...

java hadoop amazon-s3

asked aug 12 by [tddg](#)

2  
answers

#### Q: Out of memory error in Mapreduce shuffle phase

I am getting strange **errors** while running a wordcount-like mapreduce program. I have a **hadoop** cluster with 20 slaves, each having 4 GB RAM. I configured my map tasks to have a **heap** of 300MB and my ... was googling and found this link but I don't really know what to make of it: **hadoop** common link I don't understand why **hadoop** would experience any problems in copying and merging if it is able ...

hadoop mapreduce

asked oct 10 '13 by [DDW](#)

0  
answers

#### Q: FileNotFoundException while running Hadoop MR Job

and it works fine for the Sample WordCount program. I am new to **Hadoop**. I googled but could not find anything substantial. I am using **hadoop**-0.20.2-cdh3u6 on a single node setup. ...

hadoop mapreduce

asked aug 28 '13 by [Nayan](#)

2  
answers

#### Q: What is wrong with this Java for HDInsight Hadoop?

I am trying to find out why the below Java does not work when I try to run it on **hadoop**. The **error** is: c:\hadoop-training\tutorial02-jobtracker>**hadoop** jar PageStat.jar PageStat jobtra cker ... /input/visit\_5000000.txt jobtracker/output 13/07/29 11:24:50 INFO input.FileInputFormat: Total input paths to process : 1 log4j:ERROR Failed to rename [c:\Hadoop\hadoop-1.1.0-SNAPSHOT\logs ...

java hadoop hdinsight

asked jul 30 '13 by [user2244947](#)

#### A: hive collect\_set crashes query

This is probably the memory problem, since aggregates data in the memory. Try increasing **heap** size and enabling concurrent GC (via setting **Hadoop** to e.g ). This answer has more information about "GC overhead limit" **error**. ...

answered jan 14 '14 by [Nigel Tufnel](#)

2  
answers

#### Q: Hadoop: strange NullPointerException running in local mode

see in the mapper tests "processRecordsFromSuspiciousFile") but it doesn't return **errors**. I also checked through logs without any success. Is it something related to wrong or missing local-mode parameters (number threads, **heap**-memory, etc)? or something wrong in the code? ... I just started to study **Hadoop** and I have extracted an example from a book. So I have created a MapReducer to run locally that extracts temperatures from NCDC free data files. This is a sample ...

java hadoop

asked sep 22 '13 by [Randomize](#)

#### A: What is the relation between 'mapreduce.map.memory.mb' and 'mapred.map.child.java.opts' in A...

via mapred.map.child.java.opts (or mapreduce.map.java.opts in **Hadoop** 2+). If the mapper process runs out of **heap** memory, the mapper throws a java out of memory exceptions: **Error** ... physical memory

used; 970.1 MB of 1.0 GB virtual memory used. Killing container. **Hadoop** mapper is a java process and each Java process has its own **heap** memory maximum allocation settings configured ...

answered sep 20 by [user1234883](#)

2

answers

### Q: MapReduce jobs in hive-0.8.1-cdh4.0.1 Failed.

Queries in hive-0.8.1-cdh4.0.1 that invoke the Reducer results in Task Failed. The queries having MAPJOIN is working fine but JOIN gives **error**. eg: The log file shows that its due to Java **heap** space problem. ...

hive

asked oct 15 '12 by [Jickson T George](#)

1

answer

### Q: hadoop yarn single node performance tuning

(bytes) snapshot=13738057728 Virtual memory (bytes) snapshot=66425016320 Total committed **heap** usage (bytes)=10799808512 Shuffle **Errors** BAD\_ID=0 ... I have **hadoop** 2.5.2 single mode installation on my Ubuntu VM, which is: 4-core, 3GHz per core; 4G memory. This VM is not for production, only for demo and learning. Then, I wrote a vey simple map ...

performance

hadoop

mapreduce

yarn

asked dec 16 by [Tiger.Xing](#)

2

answers

### Q: setting JAVA\_HOME in linux

i set java\_home: but when i have this **error** in running **hadoop**: EDIT I set on start-all.sh but not workeed EDIT2 I edit file and set java folder but not working yet ...

java

linux

hadoop

asked dec 10 '13 by [2 8](#)

0

answers

### Q: Can I use the ResettableIterator interface of org.apache.commons.collections in my mapreduce...

and **heap** space **error** is provided here : "Getting java **heap** space **error** while running a mapreduce code for large dataset" ... into a list and then traverse the list the second time. But since I am dealing with big data the list is becoming too long and java **heap** space problem is being reported. Then I came across the Resettable ...

java

hadoop

mapreduce

asked apr 21 by [Monami Sen](#)

### A: After starting hadoop, i cannot start mahout ..!

. The **error** itself says the problem: you requested a **heap** that's too large. So, I'd go see what **heap** you requested in **Hadoop** config and check it. This is the kind of info you should post in a question. ... There's no OpenJDK 0.6 -- you mean 1.6. Yes Java 6 is supported, as you see in the project documentation. This does not seem to have anything to do with Mahout as it's an **error** from the JVM itself ...

answered apr 13 '13 by [Sean Owen](#)

### A: out of Memory Error in Hadoop

After trying so many combinations, finally I concluded the same **error** on my environment (Ubuntu 12.04, **Hadoop** 1.0.4) is due to two issues. Same as Zach Gamer mentioned above. don't forget ... to execute "ssh localhost" first. Believe or not! No ssh would throw an **error** message on Java **heap** space as well. ...

answered nov 16 '12 by [etiolap](#)

3

answers

### Q: OutOfMemory Error when running the wikipedia bayes example on mahout

i ran mahout wikipedia example with the 7 gig wiki backup.. , but when testing the classifier, i am getting the a OutOfMemory **Error** i have pasted the output below, i set the mahout **heap** size and java **heap** size to 2500m ...

java

hadoop

mahout

asked apr 9 '12 by [Nauman Bashir](#)

### A: How to insert data into Parquet table in Hive

What are the **error** messages that you get on hive server side? I had a similar problem. In the hive server log I saw some **heap** memory problems. I could solve the problem on my **hadoop** installation using higher values in mapred-site.xml ...

answered apr 26 by [woopi](#)

### A: Load snappy-compressed files into Elastic MapReduce

decompress out files. One symptom of this problem was a **heap** space **error**: When I switched to a much larger instance and cranked up the mapred.child.java.opts setting, I got a new **error**: Hadoop's snappy ... The answer is, "it can't be done." At least, not for the specific case of applying **hadoop** streaming to snappy-compressed files originating outside of **hadoop**. I (thoroughly!) explored two main ...

answered apr 4 '13 by [Abe](#)

0

answers

### Q: Can't run giraph code over Cloudera CDH

I am trying to run the shortest path code in Giraph which is installed over CDH (**Hadoop** is installed on a cluster consisting of four nodes). I tried it running like: The following was the output ...

hadoop

cloudera-cdh

giraph

asked sep 25 by [arpp](#)

0

answers

### Q: Launching jobs in a for loop

launching 50 jobs, one by one causes difficulties to **hadoop**? EDIT: Just checked the job logs, this is the first **error** occurring: And after that we see a bunch of: EDIT2: Some ideas? The **heap** ... cleared then?) EDIT3: @Donald: I haven't resized the memory for the **hadoop** daemons, so they should have a **heap** of

1GB each. The maptasks have 800 MB of **heap** from which 450 MB is used for io.sort ...

java hadoop mapreduce

asked nov 12 '13 by DDW

0

answers

### Q: Mahout error on links-simple-sorted.txt

**errors: Errors** stay the same after rerun under 2GB of **heap**. If anyone has seen it before, please let me know. Thanks! ... Following the "Mahout in Action" example of Chapter 6. links-simple-sorted.txt After format conversion ran the command of on a single node with **heap** of 1GB: After about 5 hours got the following ...

mahout

asked feb 18 '14 by user3322621

1

answer

### Q: Cause of error java.lang.IndexOutOfBoundsException: Index: 0, Size: 0

I wrote **hadoop** the program from the following Map phase As a result of execution I received the following messages I also will give Map a phase of the program which was executed successfully ... directly from a database) Help to eliminate the **error**. The **error** is in a line because probably is equal to zero. But why as a result of reading from a database (the table specified in request is completed with 6 records) wasn't filled? ...

java linux hadoop arraylist indexoutofboundsexception

asked apr 27 '13 by user2306966

### A: pig join gets OutOfMemoryError in reducer when mapred.job.shuffle.input.buffer.percent=0.70

Generally speaking, mapred.job.shuffle.input.buffer.percent=0.70 will not trigger OutOfMemory **error**, because this configuration ensures at most 70% of reducer's **heap** is used to store shuffled data ... % of **heap** in shuffle phase, which may cause OutOfMemory **error**. But in general, Pig does not have combine() in Join operator. 2) JVM manages the memory itself and divides its **heap** into Eden, S0, S1 and old ...

answered aug 14 '13 by Lijie Xu

### A: Hive: SQL request to split a table into N tables of approximately the same size?

First of all, I would dig into why you're getting **heap** size **errors**. This usually indicates a misconfigured cluster. In theory, Hive/**Hadoop** should be able to do almost everything by streaming to/from ...

answered may 17 by Joe K

1

answer

### Q: Cloudera Manager Embedded PostgreSQL Hive Metastore Server OutOfMemoryError Issue

I looked at the logs for the hive metastore and discovered java.lang.OutOfMemoryErrors: At this point, the hive metastore gets shutdown and restarted: Now, to fix this, I've changed the max **heap** ... that the actual metastore 'database' is being run as part of my cloudera manager and I'm wondering if that PostgreSQL process is running out of memory. I looked for ways to increase the java **heap** size ...

hive cloudera hue beeswax

asked may 7 '13 by Vijay Ratnagiri

### A: After starting hadoop, i cannot start mahout ..!

It's exactly what is says in the **error** message Could not reserve enough space for object **heap** check your your **hadoop** config files: and for any properties where you have allocated memory ...

answered apr 15 '13 by Ion Cojocaru

### A: Hadoop on Amazon Cloud

Have you tried Amazon Elastic MapReduce? This is a simple API that brings up **Hadoop** clusters of a specified size on demand. That's easier then to create own cluster manually. But once the jobflow ... permanently (e.g. if you are running HBase on top of **Hadoop**) you may actually need own cluster on EC2. In this case you may find Cloudera's distribution of **Hadoop** for Amazon EC2 useful. Altering **Hadoop** ...

answered nov 8 '11 by Alexey Tigarev

1

answer

### Q: How to work with Mahout?

I have tried to work with mahout using the link: http://girlincomputerscience.blogspot.in/2010/11/apache-mahout.html When i execute the command it shows: What may be the possible reasons for the **errors**? ...

apache hadoop cloud mahout recommendation-engine

asked jul 17 by user2799441

1

answer

### Q: Unable to start Hbase Thrift server

the terminal was still hanged. The page did not open but the terminal resumed showing the **error** message # # java.lang.OutOfMemoryError: Java **heap** space # -XX:OnOutOfMemoryError="kill -9 %p" # Executing /bin/sh -c "kill -9 18526"... Killed Please suggest me what to do. Regards ... I have install **Hadoop** 1.2.1 and Hbase 0.94.8 in psudo-distributed mode in my ubuntu system running Ubuntu 12.04. Everything is working fine and all the deamons are running. In order to access hbase ...

php ubuntu hadoop hbase thrift

asked feb 19 '14 by xavii

1

answer

### Q: Display the imported table in HDFS using sqoop

i have set up a single node **hadoop** cluster and configured it to work with apache Hive, now when i imported a mysql table using the following command (using sqoop) sqoop import --connect ... run the import command again it gives me the **error** that dir already exists please help me find the solution EDIT: i havent created any tables prior to the import, do i have to create a table ...

hadoop fedora hive sqoop

asked aug 4 '12 by John x

### A: Loading gz file with pig and storing it into HBase with COMPRESSION => 'GZ'

at the file extension to see what type of compression codec it was used to compress and calls the relevant decompressor. **Hadoop** first checks all the codecs installed and will report an **error** if it can't find ... the codec required for your file. So, . And this is the reason of **heap** size **error** in some mapreduce job as mapred child can't get enough space to keep the uncompressed data. . Also, there are only 3 ...

answered may 8 by [Chandra kant](#)

1  
answer

### Q: Determination of execution of mapper

on which mapper is executed. Here I give a code But despite it on one of the subordinate nodes the /export/**hadoop**-1.0.1/bin/readwaysfromdatabase.txt file wasn't created. Whether follows from here, what ...

[java](#) [linux](#) [hadoop](#) [mapreduce](#) [mapper](#)

asked apr 22 '13 by [user2306966](#)

### A: how to test Hadoop mapreduce

into CPU and **heap** dumps of **Hadoop** job. It also provides an in-depth graphical view of the MapReduce phases – Map, Reduce, Sort, Shuffle, Setup, and Cleanup. HDFS data validation Provides a simple ... and easy-to-use utility to find discrepancies and **errors** in HDFS data. User may check for data violations under various categories like data type, null values, or regular expression values. Cluster ...

answered aug 24 by [Mash](#)

1  
answer

### Q: Sqoop Oracle Hive Error

I am having a problem with the sqoop finding my table's upper bound value. It has to be 203 but its coming null as a result of which the query (in bolded ) is wrong.

[oracle](#) [hadoop](#) [hive](#) [sqoop](#)

asked apr 15 '13 by [Arjun Harish](#)

2  
answers

### Q: RHadoop reduce job failed

I am following RHadoop tutorial, <https://github.com/RevolutionAnalytics/rmr2/blob/master/docs/tutorial.md> and running the second example, but I am getting **errors** which I can't resolve. The code ... is as the following: The map job is successful, but reduce job failed, part of the **error** message is as the following: Job failed as tasks failed. failedMaps:0 failedReduces:1 could somebody help? I couldn't proceed further from here. Thanks. ...

[hadoop](#) [rhadoop](#)

asked jul 24 by [user3714023](#)

1  
answer

### Q: Reason and java.lang.NullPointerException bug fixing

I wrote the program from the following phase And as a result of execution on a cluster I received the following messages The **error** contains in a line In what cause of **error**? What it generated (owing to what it arose)? Help to remove it. ...

[java](#) [linux](#) [unix](#) [hadoop](#) [nullpointerexception](#)

asked may 1 '13 by [user2328488](#)

1  
answer

### Q: Apache Spark - MLlib - K-Means Input format

else. I use Spark 1.0.0 and my input textile (400MB) looks like this: I am not sure if I can use the first two, because in the MLlib example file there only use floats. So I also tried the last two: My **error** code in both cases is here: What am I missing? ...

[scala](#) [k-means](#) [apache-spark](#) [mllib](#)

asked aug 7 by [user3400996](#)

1  
answer

### Q: Memory issues when running Spark job on relatively large input

following **errors** in the logs: and also: How do I go about debugging such an issue? EDIT: I Found the root cause of the problem. It is this piece of code: Particularly allocating a buffer of size ...

[hadoop](#) [apache-spark](#) [google-hadoop](#)

asked oct 22 by [Yaniv Donenfeld](#)