

1
answer

Q: OOM exception in Hadoop Reduce child

I am getting **OOM** exception (Java heap space) for reduce child. In the reducer, I am appending all the values to a StringBuilder which would be the output of the reducer process. The number of values ... removed it later on. Some sample sizes of iterator are as follows: 238695, 1, 13, 673, 1, 1 etc. These are not very large values. Why do I keep getting the **OOM** exception? Any help would be valuable to me. Stack trace ...

hadoop mapreduce out-of-memory

asked oct 11 '12 by Raghava

1
answer

Q: setting hadoop job configuration programmatically

I am getting **OOM** exception (Java heap space) for reduce child. I read in the documentation that increasing the value of to -Xmx512M or more would help. Since I am not the admin, I cannot change ... that value in mapred-site.xml. I would like to set that value only for my job through the java program. I tried setting it using class as follows, but that didn't work. The version of **Hadoop** is 1.0.3 What is the proper way of setting the configuration values programmatically? ...

configuration hadoop mapreduce

asked oct 10 '12 by Raghava

2
answers

Q: Hadoop: Using cgroups for TaskTracker tasks

Is it possible to configure or **Hadoop** in a way that each process that is spawned by the TaskTracker is assigned to a specific cgroup? I want to enforce memory limits using . It is possible ... to reserve 1,5GB for the DataNode and system utilities and let the **Hadoop** TaskTracker use 6,5GB of memory. Now I start a Job using the streaming API at spawns 4 mappers and 2 reducers (each of these could ...

linux hadoop linux-kernel cgroups

asked feb 27 '13 by mt_

4
answers

Q: "hadoop namenode -format" returns a java.net.UnknownHostException

I'm currently learning **hadoop** and I'm trying to setup a single node test as defined in http://hadoop.apache.org/common/docs/current/single_node_setup.html I've configured ssh (I can log without ... a password). My server is on our intranet, behind a proxy. When I'm trying to run bin/**hadoop** namenode -format I get the following java.net.UnknownHostException exception: After that, **hadoop** ...

java configuration hadoop hdfs

asked jun 10 '11 by Pierre

1
answer

Q: hadoop single node installation - formatting error - Content is not allowed in prolog

Im trying to do a **hadoop** single node installation. I am following this article for **hadoop** installation instructions. One of the steps I need to do is to format the **hadoop** file system using ... the command hduser@dbserver:~/**hadoop**/conf\$ /home/hduser/**hadoop**/bin/**hadoop** namenode -format i get the above error message when trying to format Any suggestions to fix this issue are welcome :) . the core-site.xml file is here for your reference @zensen ~ ...

hadoop

asked feb 6 '12 by user445

1
answer

Q: Getting Fatal Error while running bin/hadoop namenode -format

Getting Fatal Error while running bin/**hadoop** namenode -format Using Windows 7 operating system, Under C:\cygwin\usr\local**hadoop**-0.20.203.0\conf edited **hadoop**-env.sh file, I have my Java class ...

java xml hadoop xmldocument

asked feb 24 '12 by Andy

A: -XX:OnOutOfMemoryError="kill -9 %p" Problem

Running as a **hadoop** option I run in to the same issues. This was the answer: Here is stdout on **OOM**: I also tried: It started, but on **OOM** it But STDERR has: sh: kill -9 1164: command not found These Wouldn't even start: ...

answered may 23 '13 by Jeremy

1
answer

Q: enabling a hadoop scheduler(Resource-aware Adaptive Scheduler)

I want to enable adaptivescheduler in **hadoop** 0.20.203.0. I have a jar file from this scheduler. (I am sure that this jar file works properly). I put jar file in HADOOP_HOME/lib and I set ... HADOOP_CLASSPATH in **hadoop**-env.sh. and I set the required property of scheduler in mapred-site.xml. when I run my cluster, all of jobtracker, datanode,.... start. But when I go to UI of scheduler(), I face ...

java hadoop

asked jul 27 '13 by Hamedreza Berenjian

0
answers

Q: Is it possible to intervene if a task fails?

I have a mapreduce job running on many urls and parsing them. I need way to handle a scenario in which one parsing task crashes on a fatal error like **OOM** error. In the normal **hadoop** behaviour a task ...

hadoop

asked may 8 '12 by user1251654

3
answers

Q: Limit CPU / Stack for Java method call?

I am using an NLP library (Stanford NER) that throws **OOM** errors for rare input documents. I plan to eventually isolate these documents and figure out what about them causes the errors ... , but this is hard to do (I'm running in **Hadoop**, so I just know the error occurs 17% through split 379/500 or something like that). As an interim solution, I'd like to be able to apply a CPU and memory limit ...

java nlp stanford-nlp

asked jul 4 '09 by Kevin Peterson

9

answers

Q: No data nodes are started

I am trying to set up **Hadoop** version 0.20.203.0 in a pseudo distributed configuration using the following guide: <http://www.javacodegeeks.com/2012/01/hadoop-modes-explained-standalone.html> After ... running the start-all.sh script I run "jps". I get this output: When I try to add information to the hdfs using: bin/**hadoop** fs -put conf input I get a error: I am not totally sure but I ...

hadoop

hdfs

asked apr 10 '12 by [Aaron S](#)

A: How to increment the hadoop counter number

will go **OOM** with such high counter values. And this is the reason behind keeping the number of counters, that can be created per Job, limited by **Hadoop**. ...

answered jun 27 '13 by [Tariq](#)

6

answers

Q: no namenode error in pseudo-mode

I'm new to **hadoop** and is in learning phase. As per **Hadoop** Definitve guide, i have set up my **hadoop** in pseudo distributed mode and everything was working fine. I was even able to execute all ...

hadoop

asked jun 22 '11 by [Anshu Basia](#)

A: HBase distributed scanner

and removing items from that queue otherwise you might **OOM** the client in case you had too many rows coming in too fast. You will also need to use concurrent structures to avoid synchronization delays. You ... can retrieve the region information using getRegionLocations on an HTable:
[http://hbase.apache.org/apidocs/org/apache/hadoop/hbase/client/HTable.html#getRegionLocations\(\)](http://hbase.apache.org/apidocs/org/apache/hadoop/hbase/client/HTable.html#getRegionLocations())
Also keep in mind ...

answered feb 27 '13 by [mibrahim](#)